

Lecture 8:

Rasterization

EEC 277, Graphics Architecture
John Owens, UC Davis, Winter 2017

No lecture Feb 14

- I'm a member of the dean's strategic planning committee!
- Feb 16: Texturing

Outline

- **Rasterization basics**
- **Pixel coverage**
- **Parameter determination**
- **Perspective correction**
- **Implementations**

What Is Rasterization?


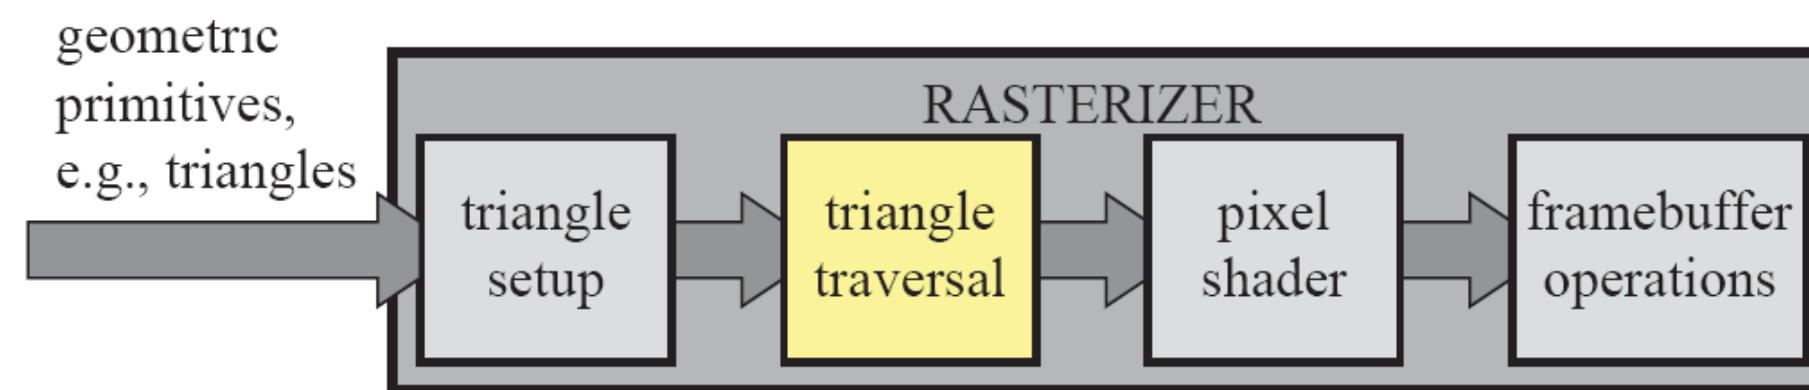
- **Input:**
 - **Screen-space geometry**
 - **Parameters per vertex**
- **Output:**
 - **Fragments**
 - **Sample (pixel) location**
 - **Parameters per fragment**
- **Rasterization has two parts:**
 - **Pixel coverage**
 - **Parameter determination**

Let's study triangle traversal


■ Critical operation in rasterizer

- without it, we have nothing
- When we have it, we will study algorithms to increase efficiency / reduce memory accesses

■ When does a pixel belong to a triangle?




Clearly, this pixel belongs to the triangle




Clearly, this pixel does NOT belong to the triangle


When does a pixel belong to a triangle?



- It all depends on where you sample!
- For (low-quality) normal rasterization, you sample in the center of the pixel:
 - If sample point is inside triangle, then pixel is inside (belongs to) the triangle.
 - In a later lecture, we will see how quality can be improved by using more than one sample per pixel




Pixel Coverage





Representing sub-pixel coordinates

- Projected points are floating-point
- Due to limited range of the coordinates, we use fixed point math (integer)
- However, we cannot round off to nearest pixel center
- Instead use sub-pixel coordinates
 - Advantage: Simpler hardware (fixed-point not floating point)
 - Advantage: Many small tris (e.g., CAD applications) disappear
- With 2 subpixel fractional bits per x, and y, we get:





Shared Pixels





Shadow Rules

- Polygon edges are in two groups: shadowed and non-shadowed
- Facing right/left: In shadow/out of shadow
- If tie, facing up/down: In shadow/out of shadow




How about when a vertex coincides with the sampling point?


- You get the same kind of problems!



- One solution: offset the subpixel grid so that sampling points never coincides with sub-pixel grid



Complex Polygon Insidedness




Jordan Curve Theorem


- Count the number of times a ray emitted from a point intersects a boundary.

- **Odd: Inside**

- **Even: Outside**




Disconnected Fragments




Insidedness Tests: Edge Equations

- $L = ax + by + c$
- $L > 0$: in ; $L = 0$: on ; $L < 0$: out
- “on”: use shadow tests




Insidedness Tests: Edge Equations

- $L = ax + by + c$
- Incremental!
- $L(x+1, y) = L(x, y) + a$




Insidedness Tests: Barycentric Formulation

- A_0 : area of red triangle; A_1 , blue; etc.
- $B_0 = A_0/(A_0+A_1+A_2)$; etc.
- If $B_i < 0$, pixel is outside




Bounding Box Traversal

- **Good: Simple**
- **Bad: Inefficient**
- **How to tell inside/outside?**



Crow's Algorithm (Scanline)


- Find lowest y
- Determine left and right edges
- Sweep upward, outputting “spans”, until edge changes
- Replace edge and keep going
- Continue until no more edges



Example: Quadrilateral




■ Pixel coverage

- Three distinct regions
 - Loop is complex
 - Must handle 1, 2, or 3 regions
- Function of
 - Screen orientation
 - Choice of \longleftrightarrow \updownarrow spans



Example: Quadrilateral

- **“All” projected quadrilaterals are non-planar**
 - Due to discrete coordinate precision
- **What if quadrilateral is concave?**
 - Concave is complex (split spans—see example)
 - Non-planar → concave for some view
- **What if quadrilateral intersects itself?**
 - A real mess (no vertex to signal change—see example)
 - Non-planar → “bowtie” for some view



All polygons are triangles (or should be)

- **Three points define a plane**
 - **Can treat all triangles as planar**
 - **Can treat all parameter surfaces as planar**
- **Triangle is always convex**
 - **Regardless of arithmetic precision**
 - **Simple rasterization, no special cases**
- **Modern GPUs decompose n -gons to triangles**
 - **SGL switched in 1990, VGX product**
 - **Optimal quadrilateral decomposition invented**

Incremental triangle traversal

$$P_i = (X_i, Y_i)$$

$$dX_i = X_{i+1} - X_i$$

$$dY_i = Y_{i+1} - Y_i$$

$$\begin{aligned} E_i(x, y) &= (x - X_i) dY_i - (y - Y_i) dX_i \\ &= A_i x + B_i y + C_i \end{aligned}$$

$E_i(x, y) = 0$: point on edge
 > 0 : outside edge
 < 0 : inside edge


Efficient incremental update:

$$dE_i(x+1, y) = E_i(x, y) + dY_i = E_i(x, y) + A_i$$

$$dE_i(x, y+1) = E_i(x, y) + dX_i = E_i(x, y) + B_i$$

Incremental update saves computation:


Only one addition per edge, per sample test



Many traversal orders are possible: backtrack, zig-zag, Hilbert/Morton curves (locality maximizing) *project idea?*


Backtrack traversal

- Have been used on mobile devices (by a Korean research group)
- Advantage: only traverse from left to right
 - Could make for more efficient memory accesses
 - Could backtrack at a faster pace (because no mem acc)



Zigzag traversal


- Simple technique that avoids backtracking
 - Otherwise, very similar
 - Can still visit a bunch of pixels outside (see next to most bottom scanline)
 - Can be solved, but requires more `Inside ()`-testing



Tiled traversal

- General idea: divide screen space into non-overlapping tiles (a tile is $w \times h$ pixels)
 - Traverse one tile at a time, and finish visiting pixels in tile before moving to next tile. Within a tile, test all pixels in parallel.
 - Can skip tiles if appropriate: entire block not in triangle ("early out"), entire block entirely within triangle ("early in")

*8x8 tile size is
common in
desktop graphics
cards*



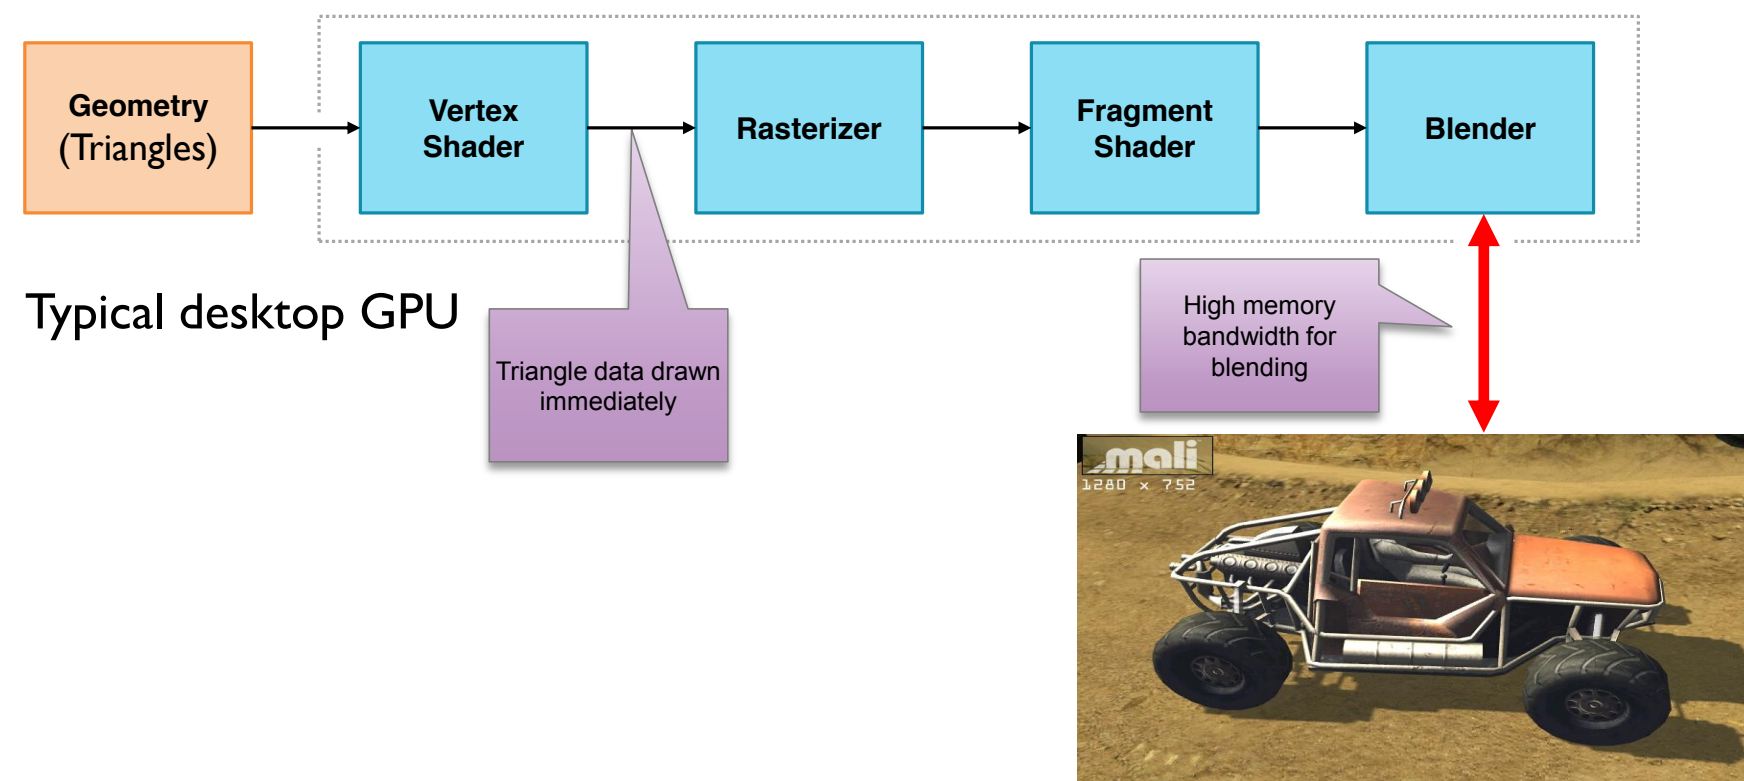
- Better because (all topics will be treated in later lectures):
 - Gives better texture cache performance
 - Enables simple culling (Zmin & Zmax)
 - Real-time buffer compression (color and depth)

Is tiled traversal that different?

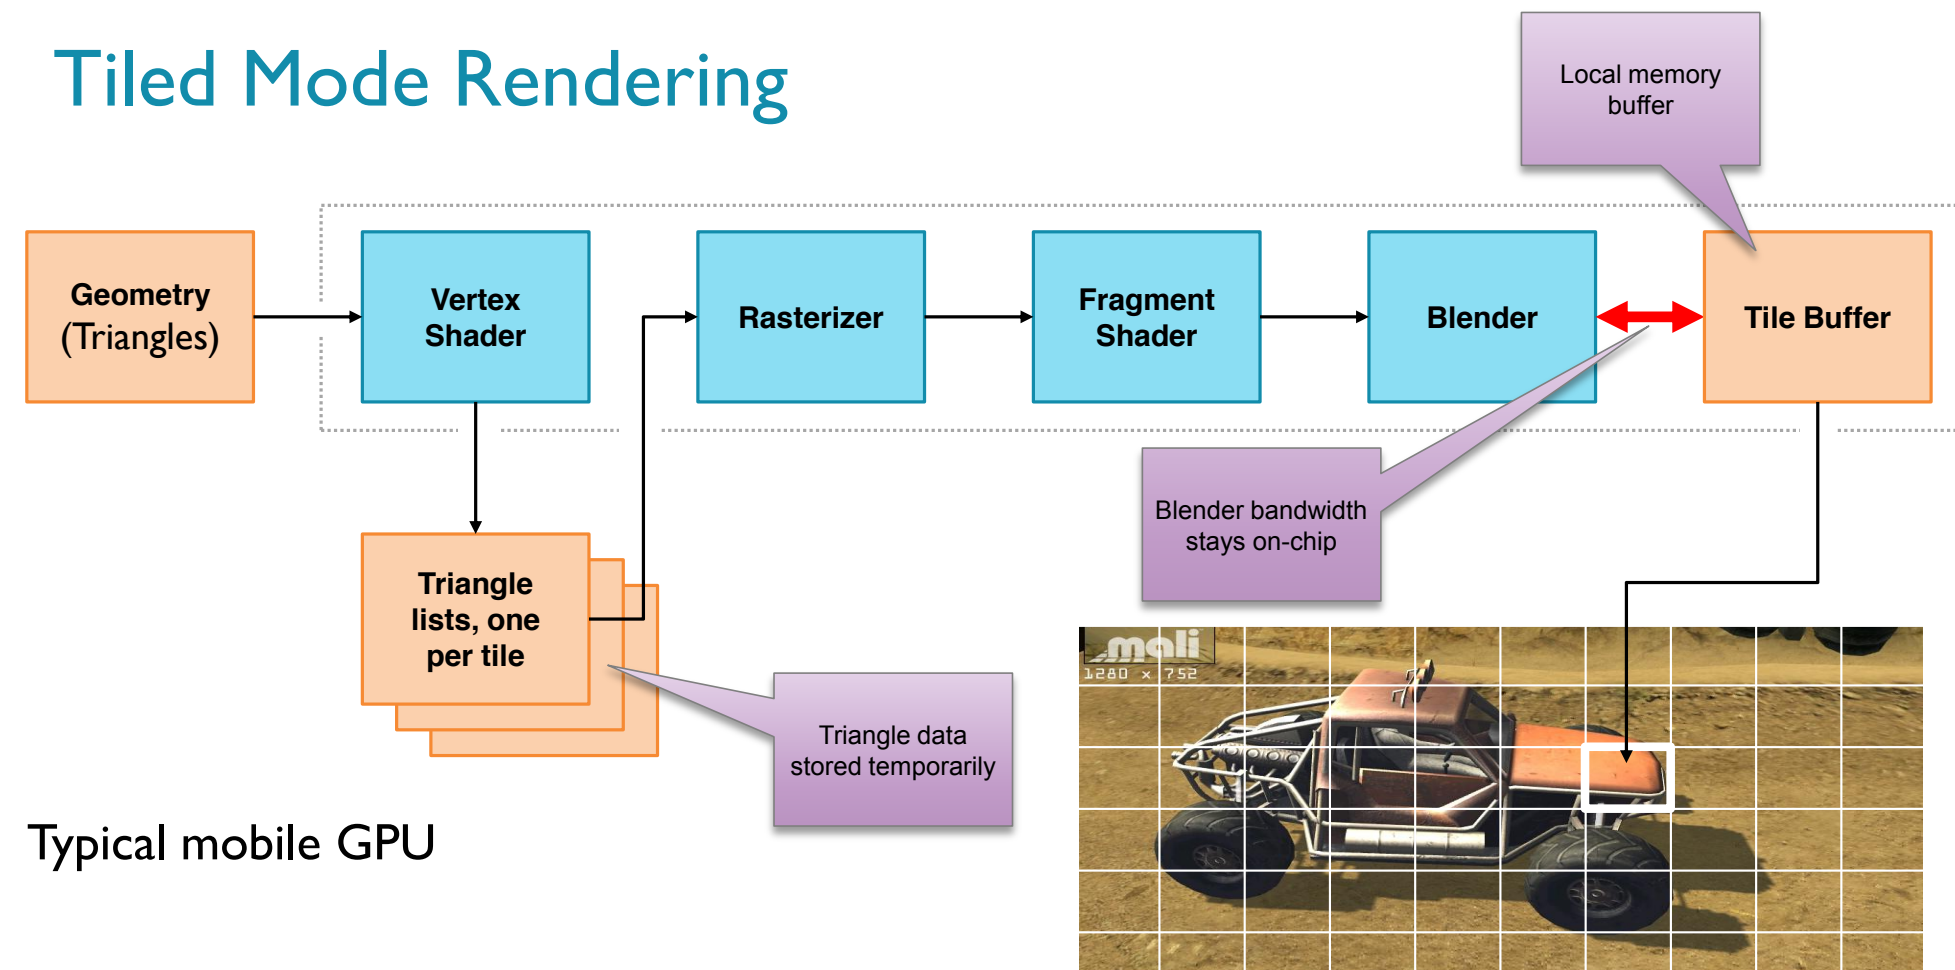
- **No, not really. We need:**
 - **I: Traverse to tiles overlapping triangle**
 - **II: Test if tile overlaps with triangle**
 - **III: Traverse pixels inside tile**
- **Previous algorithms can handle I and III**
- **II needs to be handled**
 - **Easily solved using ... edge functions!**

Tile-based rasterization

Immediate Mode Rendering



Tiled Mode Rendering



```
foreach( tile )
    foreach( primitive in tile )
        foreach( fragment in primitive in tile )
            render fragment
```

- **+ : Tile memory on-chip**
- **– : Output of vertex stage must go to main memory**
- **NVIDIA Maxwell/Pascal:**
 - Mobile-first architectures
 - Use tile-based rasterization
- **AMD, Intel: Traditional immediate-mode rasterization**

Other notes on modern rasterizers


(oral)

Outline

- Rasterization basics
- Pixel coverage
- **Parameter determination**
- Perspective correction
- Implementations

Parameter Determination

- Each triangle has several “interpolants”
- For each interpolant, what is the value at each pixel location?
- We want linear interpolation for i
- Keep in mind *setup cost* and *per-fragment cost*



Scanline Algorithm


- Big picture: Interpolation at point == interp. along edges •
interp. along scanline
- For edge e , given $(x_0, y_0), (x_1, y_1)$:
 - Calculate dx/dy and every $di/dx, di/dy$
- Start at bottom (smallest y), calculate starting x and finishing x ("span") and i_0
- Snap interpolant values to integers
- Rasterize span from x_{start} to x_{end}
 - At each step, $i += di/dx$
 - Increment $y, x += dx/dy, \text{ all } i_0 += di/dy.$
- Must handle switch in active edge for tri

Scanline Algorithm

- **Merges pixel coverage/param determination**
- **Advantages:**
 - **Efficient computationally**
 - **Uses incremental computation**
 - **Little wasted work**
 - **Good for large triangles**
- **Disadvantages**
 - **Complex control**
 - **Lots of overhead for small triangles**


Barycentric Formulation

- a_0 : area of red triangle; a_1 , blue; etc.
- $b_0 = a_0/(a_0+a_1+a_2)$; etc.
 - Have we calculated $a_0+a_1+a_2$ before?
- $i = b_0i_0 + b_1i_1 + b_2i_2$
- a_n as function of x, y ?




Barycentric Formulation

- What is area as $\text{fn}(P=(x,y))$?
- $a_0 = 1/2 |PP_1 \times PP_2|$
- $= 1/2 \begin{vmatrix} x_1-x & y_1-y \\ x_2-x & y_2-y \end{vmatrix}$
- $= (y_1-y_2)x + (x_2-x_1)y + (x_1y_2-x_2y_1)$
- $= \alpha_0 x + \beta_0 y + \gamma_0$
- Must calculate a, a_0, a_1, a_2



Barycentric Formulation

- a_0 : area of red triangle; a_1 , blue; etc.
- $b_0 = a_0/(a_0+a_1+a_2)$; etc.
- $i = b_0i_0 + b_1i_1 + b_2i_2$ ($b \cdot i$)
- a_n as function of x, y ?
- $a_0 = \alpha_0x + \beta_0y + \gamma_0$
- So areas can be calculated linearly too!




Building a Rasterizer

- **Triangle setup (work per triangle)**
- **Scanline:**
 - **Edge equations (one divide/edge)**
 - **Find smallest y / sort triangle**
- **Barycentric:**
 - **Calculate $A, \alpha_n, \beta_n, \gamma_n$**
- **Fragment generation (work per fragment)**
 - **Scanline: incremental evaluation**
 - **Barycentric: incrementally calculate b_n , dot product to find interpolants**

Pineda Rasterization

- Rasterize multiple pixels at once—deltas are all linear
- Edge equations used as masks



Juan Pineda, "A Parallel Algorithm for Polygon Rasterization", Siggraph '88

GeForce Principles II

- Try to load balance the pixel engine and the vertex engine
 - **Pixels are stamped out in 2x2 blocks per clock**
 - **Or in 2x1, or in 1x1 blocks when using heavy MultiTexture**
- Vertices are cached before TnL, after TnL and in the primitive assembly area
 - Effective use of these caches is critical for best performance




Outline

- Rasterization basics
- Pixel coverage
- Parameter determination
- **Perspective correction**
- Implementations

Parameter Determination


- Each triangle has several “interpolants”
- For each interpolant, what is the value at each pixel location?
- We want linear interpolation for i (in object space)




Perspective Correct Interpolation

- Up to this point, we've considered linear interpolation in screen space for interpolants.
- But we can't interpolate u, v linearly. Why?


Projection to straight lines




Linear Perspective



Correct Linear Perspective



Incorrect Perspective



Linear Interpolation, Bad
Perspective Interpolation, Good




Perspective Correct Interpolation

- Terminology:
- Eye space: x_e, y_e
- Screen space before divide: x_s, y_s, w_s
- Pixel coordinates: $x = x_s/w_s, y = y_s/w_s$
- Want $u, v = f(x, y)$
- Recall we used to do affine mapping of
 $u(x, y) = \alpha_0 x + \beta_0 y + \gamma_0$
- Assume affine mapping between u, v and x_e, y_e
- Could do matrix transform for each pixel
- Blackboard!

Perspective Correct Interpolation

- We can't interpolate u, v linearly.
- But we can interpolate u/w linearly and $1/w$ linearly.
- So instead of interpolating u, v , etc. ...
- We interpolate $u/w, v/w, 1/w$, etc. ...
- Requires divide per vertex for setup, divide per pixel during rasterization.
- How about rgb, z , normals, etc.?

Alternatives



Perspective Correct Interpolation

■ Scanline:

- Interpolate $1/w$ and i/w
- At each pixel, take $1/(1/w)$ and multiply for each interpolant

■ Barycentric:

- $b_0 = a_0/(a_0+a_1+a_2)$ becomes
- $b_0 = w_1w_2a_0/(w_1w_2a_0+w_2w_0a_1+w_0w_1a_2)$
- Derivatives easily computable too

Perspective Correct Interpolation

- What if w was constant across a scanline?
- (Blackboard, part 2)

Doom (first full-screen textured sw renderer)



Outline

- **Rasterization basics**
- **Pixel coverage**
- **Parameter determination**
- **Perspective correction**
- **Implementations (all slides, Kurt Akeley)**

Triangle Rasterization Examples


- **Gouraud shaded (GTX)**
- **Per-pixel evaluation (Pixel Planes 4)**
- **Edge walk, planar parameter (VGX)**
- **Barycentric direct evaluation (InfiniteReality)**
- **Small tiles (Bali—proposed)**
- **Homogeneous recursive descent (NVIDIA)**

Algorithm properties


- **Setup and execution cost**
 - **Per-triangle**
 - **Per-fragment**
- **Ability to parallelize**
- **Ability to cull to a rectangular screen region**
 - **To support tiling**
 - **To support “scissoring”**

Gouraud shaded (GTX)

- **Two stage algorithm**
 - **DDA edge walk**
 - **fragment selection**
 - **parameter assignment**
 - **DDA scan-line walk**
 - **parameter assignment only**
- **Requires expensive scan-line setup**
 - **Location of first sample is non-unit distance from edge**
- **Cannot scissor efficiently**
- **Works on quadrilaterals**



Edge walk, planar evaluation (VGX)



Edge walk, planar evaluation (VGX)


■ Hybrid algorithm

- Edge DDA walk for fragment selection
 - Efficient generation of conservative fragment set
- Sample DDA walk for parameter assignment
 - Never step off sample grid, so
 - Never have to make sub-pixel adjustment


■ Scissor cull possible

- Adds complexity to edge walk

■ Sample walk simplifies parallelism



Interpolation outside the triangle



Barycentric (InfiniteReality)

- **Hybrid algorithm**
 - **Approximate edge walk for fragment selection**
 - **Pineda edge functions used to generate AA masks**
 - **Direct barycentric evaluation for parameter assignment**
 - **Barycentric coordinates are DDA walked on grid**
 - **Minimizes setup cost**
 - **Additional computational complexity accepted**
 - **Handles small triangles well**
- **Scissor cull implemented**
 - **Supports “guard band clipping”**

Small tiles (Bali—proposed)

- **Framebuffer tiled into $n \times n$ (16x16) regions**
 - **Each tile is owned by a separate engine**
- **Two separate rasterizations**
 - **Tile selection (avoid broadcast, conservative)**
 - **Fragment selection and parameter assignment**
- **Parallelizes well**
- **Handles small triangles well**
- **Scissors well**
 - **At tile selection stage**

Homogeneous recursive descent

- **Rasterizes unprojected, unclipped geometry**
 - **Huge improvement for geometry processing!**
 - **Interpolates clip-plane distances**
- **I have heard ~2001 NVIDIA implemented this**
 - **But is not well documented**
 - **Read Olano and Greer**
 - **Parameter assignment precision has many pitfalls — I have heard NVIDIA gave up on this technique for this reason**
 - **Watch out for infinities!**
- **Recursive descent**
 - **Scissors well**
 - **Drives $n \times n$ (2×2) parallel fragment generation**
- **Cannot generate perspective-incorrect parameter values**

Next Lecture—Texturing