

Topics on VR




VR Today





What is VR?



Augmented Reality



Mixed Reality





360° Images & Video: The Jury's Still Out

- Capture a scene using an array of cameras
 - either stationary or moving
- Rendered images
 - *omni-directional Stereo* (ODS)
 - *equirectangular* projection



Device Deployment



Rendering for VR




Frame Rate

- Video games target 60 Hz refresh
 - single view, very high resolution (UHD / 4K)
 - 16.6 ms/frame
- VR is different:
 - two eyes
 - high resolution (HD / 1080)
 - need time for compositing
- We'd like 120 Hz (8.3 ms/frame)
 - settle for 90 Hz (11 ms/frame)

60 Hertz
30 Hurts



Deferred Rendering

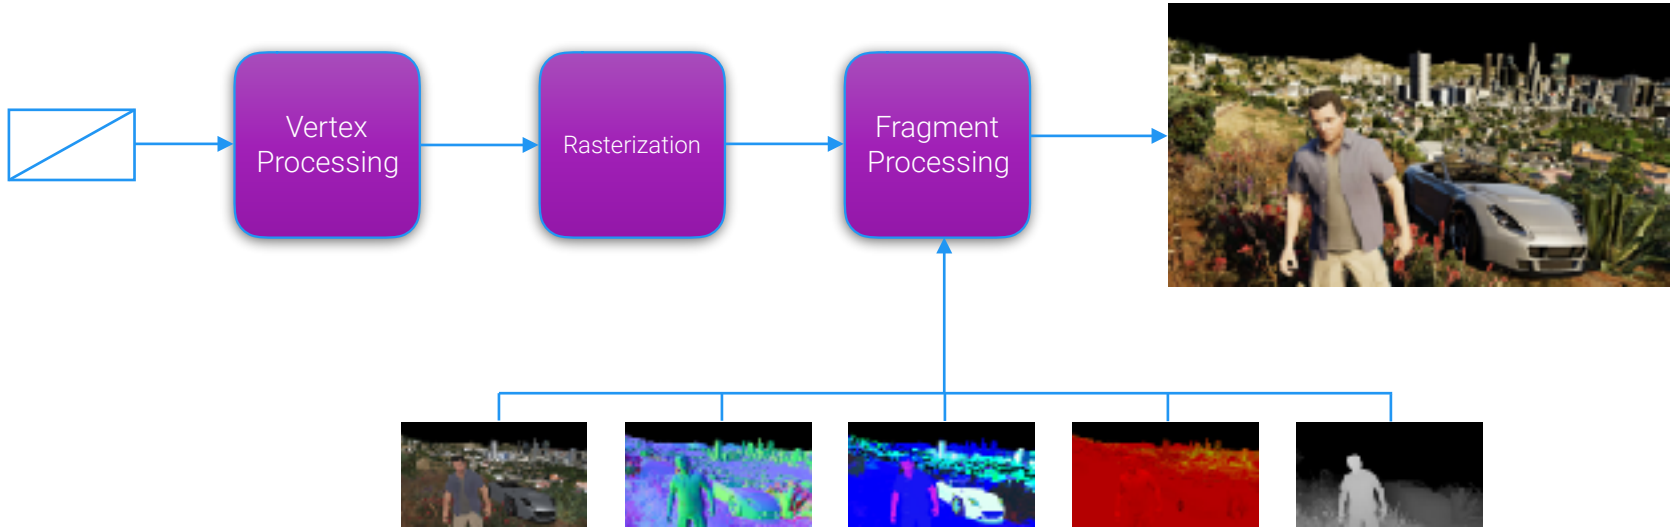


<http://www.adriancourreges.com/blog/2015/11/02/gta-v-graphics-study/>

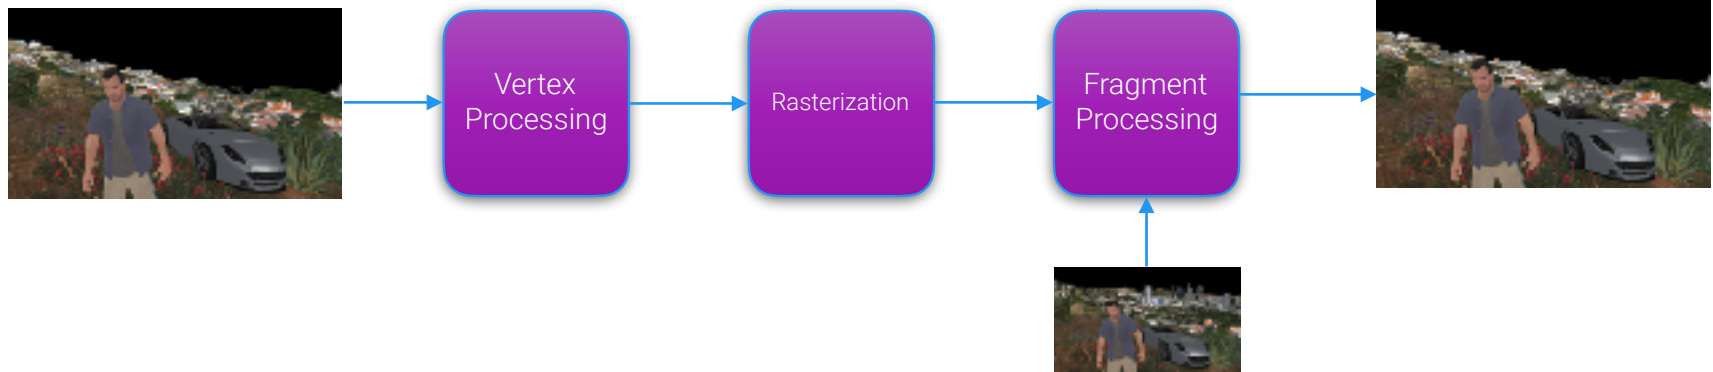


Deferred Rendering (Pass 2)

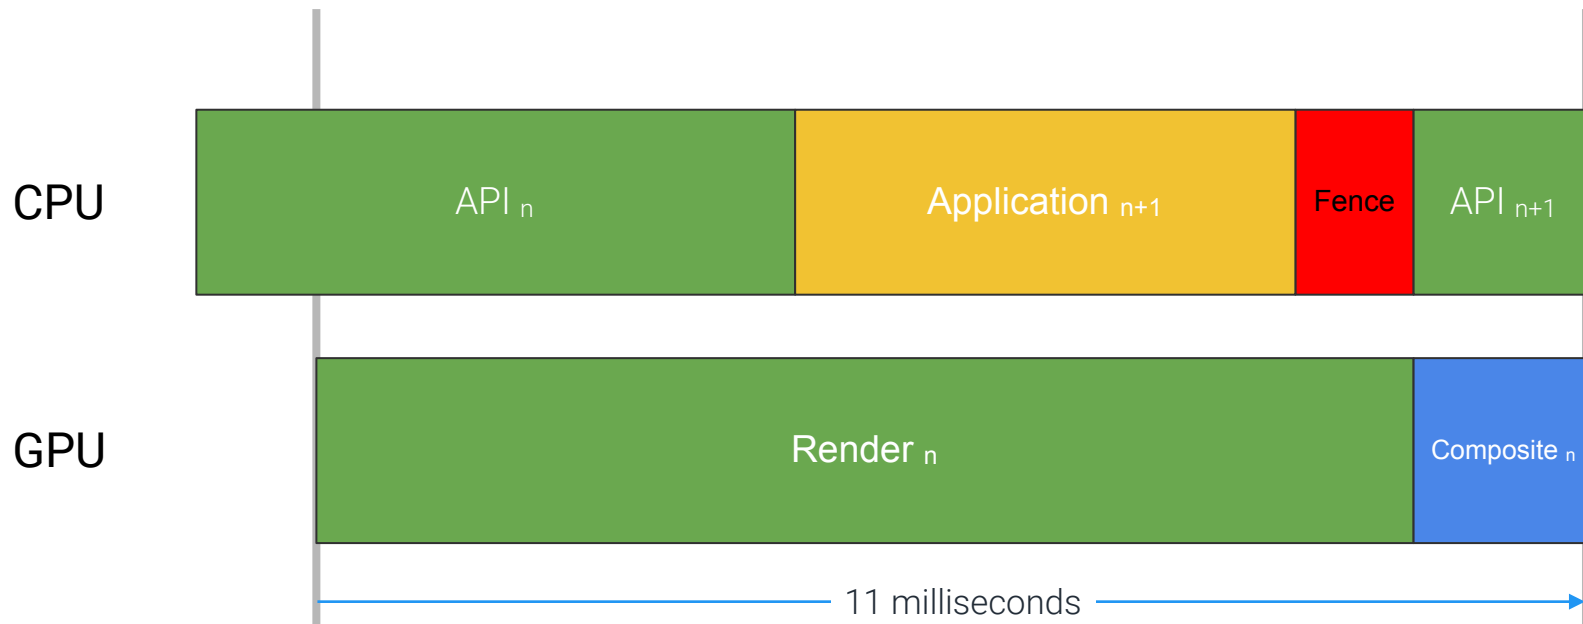
Final Image Generation



Forward Rendering




Frame Phases (@ 90Hz)



Generating Stereo Pairs

- Image for each eye generated separately
 - same frustum separated by *inter-pupillary distance* (IPD)
- API extensions to "optimize" rendering
 - application provides an array of viewing transforms
 - driver executes each draw call twice, updating index into viewing transforms



Remember, for symmetric projections: $M[0][0] = \frac{n}{f} = \frac{1}{\tan(\theta)}$

Given:

$$\frac{r}{n} = \frac{r + \frac{ipd}{2}}{n_c}$$

Solving for n_c :


$$n_c = n + \frac{ipd}{2} \frac{1}{\tan(\theta)} = n + \frac{ipd}{2} M[0][0]$$

So we only need to translate the camera back by $\frac{ipd}{2} M[0][0]$ to include the view volume of both frusta (assuming infinite far plane).



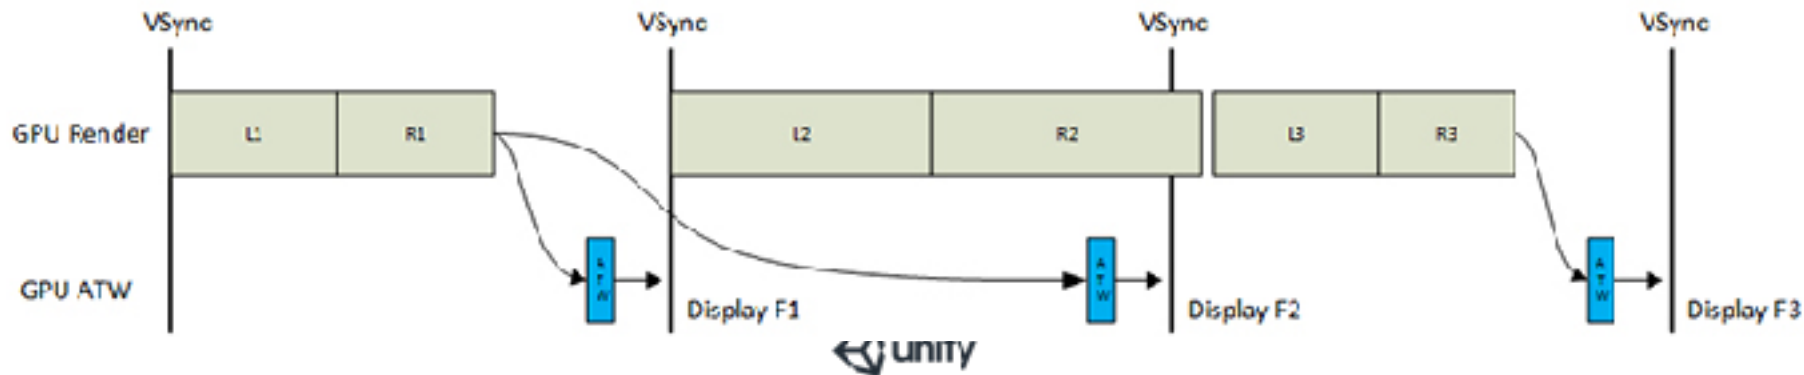
Compositing

- Device drivers include a compositing facility
- Warp each generated eye texture to match device optics



Async Time Warp

- Oculus technology for dealing with frame overruns
- Previous frame warped using updated pose information by compositor
- Helps prevent judder
- Requires GPU preemption




Triangles vs. SDFs

- Triangles aren't the (perhaps) the best solution in many cases
- *Signed Distance Fields* (SDFs) represent models as implicit surfaces
- rendered using *ray marching*
 - either rendered entirely in a compute shader
 - or full-screen quad and a helluva fragment shader



shadertoy.com



unbound.io

Input, Gestures, & Haptics



Controlling the View: DOFs

- Mice & keyboards are easy
 - 2D screen coordinates
 - binary mouse buttons & keyboard events
- Controllers provide much more data
 - *degrees of freedom* (DOFs)
 - 3-DOF
 - 6-DOF




Events vs. Gestures

- Events are easy to put together
 - usually represented as some sort of state machine
 - clear state transitions based on discrete events
- Gestures can be more difficult
 - filter input event stream
 - constrain DOFs



Pose, Gaze, & Eye Tracking

- Computer vision and deep learning techniques can identify and track objects in real time
- Tobii - commercial eye-tracking
 - integrated into several laptops
 - similar technology being built into HMDs
- Apical - real-time computer vision IP
 - identifies and tracks objects in real-time
 - returns location and pose data



Foveated Rendering (Region of Interest)

- Your eye has more resolution in its central vision area (fovea)
- render more accurately in gaze direction
 - lower levels-of-detail or resolution outside gaze cone
- Extensions or hacks required on desktop
 - amazingly simple on mobile (tiled renderers)



Current Challenges



Wires vs. Latency

- Tethered headsets are a drag
- Wireless headsets are coming
 - unclear about constraints of latency and interference
- Pack entire system into a backpack
 - [The Void](#)



Cloud vs. On-Device

- Battery-powered devices still have a ways to go with rendering performance
 - they can do pretty much everything a discrete GPU, just perhaps not as fast
- Can we render in the cloud and send to a device?
 - latency issues
 - pose information up
 - rendered framed down
 - image delivery
 - encode as H.265 (supports constant 90 Hz video encode / decode)



Thank you

